College of Science and Engineering
James Cook University
Australia
nuonan.ouyang@my.jcu.edu.au

22 October 2026

Editor-in-Chief and Associate Editor
IEEE Transactions on Sustainable Computing

Dear Editor-in-Chief and Associate Editor,

Please find enclosed the major revision of manuscript TSUSC-2026-07-0200, "Shielded Q-learning for Thermally-Safe, Energy-Aware Multi-Model Intrusion Detection Scheduling on Lightweight Edge Devices."

The revision responds to the requests for simpler baselines, paired statistics, online execution, controller overhead, physical energy measurement, guard coverage, and a second hardware configuration. The original execution artifacts were unavailable, so the empirical pipeline was reconstructed; the revised tables are replacement evidence, not a claim of numerical reproduction of the submitted replay results.

During implementation review we identified inconsistencies in the earlier Pi 4B online worker and its energy accounting. Those results are excluded. The enclosed manuscript uses the corrected, configuration-bound R6 campaign: 60 accepted online runs, 36,000 actual 16-record windows, and contemporaneous KM003C power traces. Technical failures and their replacements remain archived. The paper defines the timing and electrical boundaries and reports false-positive rates, per-seed behavior, and the limitations of the available load telemetry.

The corrected evidence does not establish a practical advantage for the Q-learning layer. Safe-Greedy performs better on the Pi 3B+ binding workload; on Pi 4B the two methods are close under lower loads, while FSSQL-R degrades under the highest load. A post hoc state-support analysis identifies unvisited Q states and a fixed default-action path. We do not claim universal shielding necessity, full-state empirical convergence, or consistent physical energy savings.

We also address the submitted statement that test-set detector statistics supplied the reward. The revised objective uses validation statistics only. The saved official detector test was evaluated on 11 September 2026. October implementation corrections and Q-policy reconstruction occurred afterwards using development data without reopening the test; detector, preprocessing, threshold, and validation-quality hashes remain unchanged. This chronology and the limits of historical provenance are stated explicitly.

A point-by-point response and supplementary material accompany the revised main manuscript. Thank you for considering the revision.

Sincerely,

Nuonan Ouyang
Corresponding Author